

7 On-policy Value Approximation

Melih Kandemir

Özyeğin University
Computer Science Department
melih.kandemir@ozyegin.edu.tr

7 Nov 2017

Value function approximation

Up till now we have stored the value functions in tabular form, but what if

- ▶ the state space is too large to allocate memory for every state?
- ▶ or, the state description is very high-dimensional (e.g. a Go table, or the scene image)?

The remedy is to approximate the value function v_π by a function \hat{v}_π parametrized by a vector \mathbf{w}

$$\hat{v}(s; \mathbf{w}) \approx v_\pi(s)$$

that applies its parameters $\mathbf{w} \in \mathbb{R}^d$ on a d -dimensional vector of state features.

Value function approximation as supervised learning problem

- ▶ Typically, $d \ll |\mathcal{S}|$, where $|\mathcal{S}|$ is the state count.
- ▶ Hence, unlike the tabular approach, updating a single state affects the values of many other states!
- ▶ This attempt to *generalize* the value function over states is essentially the well-known *regression* setting in classical *supervised* learning.

RL-friendly supervised learning models

- ▶ **Online:** In RL, samples are generated on the fly, in a non-i.i.d. fashion. Hence, the chosen learning algorithm should be eligible for online learning (train incrementally by introducing one sample at a time).
- ▶ **Non-stationary:** The behavior of most RL target functions is *non-stationary* (i.e. the data distribution changes over time). The supervised learning method should adaptively account for drifts in data distributions.

Stochastic Gradient Descent (SGD)

Given a data set $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$ of N input observations \mathbf{x}_n , the corresponding output observations y_n , and an estimator $f(\mathbf{x}_n, \mathbf{w})$, in a typical regression task, we aim to minimize the Mean Squared Error (MSE)

$$MSE = \frac{1}{N} \sum_{n=1}^N \left[y_n - f(\mathbf{x}_n, \mathbf{w}_t) \right]^2$$

subject to \mathbf{w} by gradient descent

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \frac{1}{N} \sum_{n=1}^N \left[y_n - f(\mathbf{x}_n, \mathbf{w}_t) \right] \nabla f(\mathbf{x}_n, \mathbf{w}_t).$$

Stochastic Gradient Descent (SGD)

One update requires a full pass on the entire data set

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \frac{1}{N} \sum_{n=1}^N \left[y_n - f(\mathbf{x}_n, \mathbf{w}_t) \right] \nabla f(\mathbf{x}_n, \mathbf{w}_t).$$

This

- ▶ is expensive if the data set is large.
- ▶ delays model fit.

Try out the divide-and-conquer strategy instead!

Stochastic Gradient Descent (SGD)

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \underbrace{\frac{1}{N} \sum_{n=1}^N \left[y_n - f(\mathbf{x}_n, \mathbf{w}_t) \right] \nabla f(\mathbf{x}_n, \mathbf{w}_t)}_{\text{true gradient}}.$$

Approximate the true gradient by sampling a random data point $\{\mathbf{x}_j, y_j\}$ and assuming that \mathcal{D} consists of N copies of it. Then the gradient-descent update reads

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \underbrace{\left[y_j - f(\mathbf{x}_j, \mathbf{w}_t) \right] \nabla f(\mathbf{x}_j, \mathbf{w}_t)}_{\text{stochastic gradient}}.$$

Stochastic Gradient Descent (SGD)

SGD can be generalized to updating by small randomly-chosen data subsets, called **minibatches**, trivially. Given a random subset $\{(\mathbf{x}_{i(1)}, y_{i(1)}), \dots, (\mathbf{x}_{i(\tilde{N})}, y_{i(\tilde{N})})\}$ of size $\tilde{N} \ll N$, where $i(n)$ is a random permutation, perform the update

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \frac{1}{\tilde{N}} \sum_{n=1}^{\tilde{N}} \left[y_{i(n)} - f(\mathbf{x}_{i(n)}, \mathbf{w}_t) \right] \nabla f(\mathbf{x}_{i(n)}, \mathbf{w}_t).$$

Robbins-Monro Theorem (1951)

The stochastic gradient is an *unbiased* estimate of the true gradient if updates are performed following a learning rate series satisfying the two properties below

$$\sum_{t=1}^{\infty} \epsilon_t = \infty, \quad (1)$$

$$\sum_{t=1}^{\infty} \epsilon_t^2 < \infty. \quad (2)$$

- (1) reach at points arbitrarily far away
- (2) stop learning at some point

Data Sets in RL

Monte Carlo

$$\mathcal{D}_{MC} = \begin{bmatrix} (\phi(s_1), G_1), \\ (\phi(s_2), G_2), \\ \vdots \\ (\phi(s_T), G_T) \end{bmatrix}$$

TD(0)

$$\mathcal{D}_{TD} = \begin{bmatrix} (\phi(s_1), R_1 + \gamma \hat{v}(s_2, \mathbf{w})), \\ (\phi(s_2), R_2 + \gamma \hat{v}(s_3, \mathbf{w})), \\ \vdots \\ (\phi(s_T), R_T) \end{bmatrix}$$

Legend: (input,output)

SGD with and without replacement

- ▶ **without replacement:**

- ▶ Choose a sample only once until all samples are chosen (i.e. until an **epoch** is complete).
- ▶ More common in standard ML due mainly to practical reasons (to decide whether the epoch is over).

- ▶ **with replacement:**

- ▶ Keep random sampling without caring about coverage.
- ▶ Allow multiple selection of a sample within an epoch.
- ▶ More common in RL.
- ▶ Called **Experience Replay**.

The Mean Squared Value Error (MSVE)

$$MSVE(\mathbf{w}) \triangleq \sum_{s \in \mathcal{S}} \mu(s) \left[v_{\pi}(s) - \hat{v}(s, \mathbf{w}) \right]^2$$

where $\mu(s) \geq 0$ is a distribution assigning each state a degree of importance.

Here,

- ▶ $v_{\pi}(s)$ is the true value of state s , which is our *target*,
- ▶ and $\hat{v}(s, \mathbf{w})$ is a prediction estimating this target.

If we choose $\mu(s)$ to be the fraction of time spent in s , then we perform *on-policy RL*.

SGD for value approximation

$$\begin{aligned}\mathbf{w}_{t+1} &\leftarrow \mathbf{w}_t - \frac{1}{2}\alpha \nabla \left[v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t) \right]^2 \\ &\leftarrow \mathbf{w}_t + \alpha \left[v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t) \right] \nabla \hat{v}(S_t, \mathbf{w}_t) \\ &\leftarrow \mathbf{w}_t + \alpha \left[U_t - \hat{v}(S_t, \mathbf{w}_t) \right] \nabla \hat{v}(S_t, \mathbf{w}_t),\end{aligned}$$

where U_t is the approximation target for $v_\pi(S_t)$. For Monte Carlo RL, $U_t = G_t$.

Gradient MC for value approximation

Input: $\pi, \hat{v} : \mathcal{S} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Initialize \mathbf{w}

repeat forever

 Generate an episode $S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_t$ from π

 For $t = 0, 1, \dots, T - 1$:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left[G_t - \hat{v}(S_t, \mathbf{w}) \right] \nabla \hat{v}(S_t, \mathbf{w})$$

Convergence of Gradient Monte Carlo Value Approximation

- ▶ If $\mathbb{E}[U_t] = v_\pi(S_t)$, then Robbins-Monro is satisfied following a proper learning rate series ϵ_t . Hence, $\hat{v}(S_t, \mathbf{w})$ converges to a local optimum.
- ▶ $G_t = \mathbb{E}[U_t] = v_\pi(S_t)$ by definition.

Semi-gradient Methods

The situation about convergence is more challenging for bootstrapping methods, as the target itself also depends on the parameters being learned!

For the TD(0) target, the gradient-descent update reads

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \left[\underbrace{R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w})}_{U_t \text{ for TD(0)}} - \hat{v}(S_t, \mathbf{w}_t) \right] \nabla \hat{v}(S_t, \mathbf{w}_t).$$

- ▶ Robbins-Monro no longer applies!
- ▶ Accounts for the effect of the change on the estimate by updating \mathbf{w} , but ignores its effect on the target.
- ▶ Called **semi-gradient methods**.

Semi-gradient TD(0) for prediction

Input: $\pi, \hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$ s.t. $\hat{v}(\text{terminal}, \cdot) = 0$.

Initialize \mathbf{w}

repeat (for each episode)

Initialize S

repeat (for each step of episode)

Choose $A \sim \pi(\cdot|S)$

Take action A , observe R, S'

$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left[R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w}) \right] \nabla \hat{v}(S, \mathbf{w})$

$S \leftarrow S'$

until S' is terminal

Value approximation with a linear estimator

For a linear value estimator $\hat{v}(S_t, \mathbf{w}) = \phi_t^T \mathbf{w}$, where $\phi_t \in \mathbb{R}^d$ is a d -dimensional feature vector for state S_t , the TD(0) semi-gradient update rule is

$$\begin{aligned}\mathbf{w}_{t+1} &\leftarrow \mathbf{w}_t + \alpha \left[R_{t+1} + \gamma \phi_{t+1}^T \mathbf{w}_t - \phi_t^T \mathbf{w}_t \right] \phi_t \\ &\quad \mathbf{w}_t + \alpha \left[R_{t+1} \phi_t - \phi_t (\phi_t - \gamma \phi_{t+1})^T \mathbf{w}_t \right].\end{aligned}$$

At the steady-state, the expected update reads

$$\mathbb{E}[\mathbf{w}_{t+1} | \mathbf{w}_t] = \mathbf{w}_t + \alpha (\mathbf{b} - \mathbf{A} \mathbf{w}_t),$$

where

$$\mathbf{b} = \mathbb{E}[R_{t+1} \phi_t] \in \mathbb{R}^d, \quad \mathbf{A} = \mathbb{E} \left[\phi_t (\phi_t - \gamma \phi_{t+1})^T \right] \in \mathbb{R}^{d \times d}.$$

The linear system converges to \mathbf{w}_{TD} (called the **TD fixpoint**) if the update results in no change

$$\mathbf{b} - \mathbf{A} \mathbf{w}_{TD} = \mathbf{0} \Rightarrow \mathbf{w}_{TD} = \mathbf{A}^{-1} \mathbf{b}.$$

TD error bound for the linear estimator

The linear value estimator of TD satisfies the error bound below

$$\min_{\mathbf{w}} MSVE(\mathbf{w}_{TD}) \leq \frac{1}{1-\gamma} \min_{\mathbf{w}} MSVE(\mathbf{w}),$$

which is not such a tight estimator for $\gamma \approx 1$.

Least-squares TD (LSTD)

Actually, why should we do gradient-descent for the linear model? We know from linear regression that it has analytical solution

$$\mathbf{w}_{t+1} = \hat{\mathbf{A}}_t^{-1} \hat{\mathbf{b}}_t,$$

where

$$\hat{\mathbf{A}}_t = \sum_{k=0}^t \phi_k (\phi_k - \gamma \phi_{k+1})^T + \epsilon \mathbf{I},$$
$$\hat{\mathbf{b}}_t = \sum_{k=0}^t R_{t+1} \phi_k$$

for some $0 < \epsilon \ll 1$ that guarantees invertibility.

Online learning with LSTD

The Sherman-Morrison formula allows incremental calculation of the matrix inverse

$$(\mathbf{M} + \mathbf{u}\mathbf{v}^T)^{-1} = \mathbf{M}^{-1} - \frac{\mathbf{M}^{-1}\mathbf{u}\mathbf{v}^T\mathbf{M}^{-1}}{1 + \mathbf{v}^T\mathbf{M}^{-1}\mathbf{u}}.$$

Applied to LSTD, we get

$$\begin{aligned}\hat{\mathbf{A}}_t^{-1} &= \left(\underbrace{\hat{\mathbf{A}}_{t-1}}_{\mathbf{M}} + \underbrace{\phi_t}_{\mathbf{u}} \underbrace{(\phi_t - \gamma\phi_{t+1})^T}_{\mathbf{v}} \right)^{-1} \\ &= \hat{\mathbf{A}}_{t-1}^{-1} - \frac{\hat{\mathbf{A}}_{t-1}^{-1}\phi_t(\phi_t - \gamma\phi_{t+1})^T\hat{\mathbf{A}}_{t-1}^{-1}}{1 + (\phi_t - \gamma\phi_{t+1})^T\hat{\mathbf{A}}_{t-1}^{-1}\phi_t}.\end{aligned}$$

The LSTD algorithm

Input: Feature representation $\phi(s) \in \mathbb{R}^d$, $\phi(\text{terminal}) = \mathbf{0}$

$\hat{\mathbf{A}}^{-1} \leftarrow \epsilon^{-1} \mathbf{I}$

$\mathbf{b} \leftarrow \mathbf{0}$

repeat (for each episode)

Initialize S , obtain ϕ

repeat (for each step of the episode)

Choose $A \sim \pi(\cdot|S)$

Take action A , observe R, S' , obtain ϕ'

$\mathbf{v} \leftarrow \hat{\mathbf{A}}^{-T}(\phi - \gamma\phi')$

$\hat{\mathbf{A}}^{-1} \leftarrow \hat{\mathbf{A}}^{-1} - (\hat{\mathbf{A}}^{-1}\phi)\mathbf{v}^T / (1 + \mathbf{v}^T\phi)$

$\hat{\mathbf{b}} \leftarrow \hat{\mathbf{b}} + R\phi$

$\mathbf{w} \leftarrow \hat{\mathbf{A}}^{-1}\hat{\mathbf{b}}$

$S \leftarrow S', \phi \leftarrow \phi'$

until S' is terminal

On-policy control with approximation

As was before, we need to move from $v(s)$ to $q(s, a)$

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \left[U_t - \hat{q}(S_t, A_t, \mathbf{w}_t) \right] \nabla \hat{q}(S_t, A_t, \mathbf{w}_t).$$

The TD(0) gradient-descent update is

$$\begin{aligned} \mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \alpha \bigg[&R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t) \\ &- \hat{q}(S_t, A_t, \mathbf{w}_t) \bigg] \nabla \hat{q}(S_t, A_t, \mathbf{w}_t). \end{aligned}$$

Episodic semi-gradient Sarsa for control

Input: a differentiable $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$

Initialize $\mathbf{w} \in \mathbb{R}^d$

repeat (for each episode)

$S, A \leftarrow$ initial state and action for episode (e.g. ϵ -greedy)

repeat (for each step of episode)

Take action A , observe R, S'

If S' is terminal

$\mathbf{w} \leftarrow \mathbf{w} + \alpha[R - \hat{q}(S, A, \mathbf{w})]\nabla\hat{q}(S, A, \mathbf{w})$

Go to next episode

Choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g. ϵ -greedy)

$\mathbf{w} \leftarrow \mathbf{w} + \alpha[R + \gamma\hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})]\nabla\hat{q}(S, A, \mathbf{w})$

$S \leftarrow S'; A \leftarrow A'$

Example: Mountain Car

- ▶ Drive a car out of a U-shaped valley.
- ▶ Gravity is stronger than the car's engine.
- ▶ **Cost-to-go function:** $-\max_a \hat{q}(s, a, \mathbf{w})$.
- ▶ **Reward:** -1 per time step, +100 for reaching the goal.
- ▶ **Actions:**
 - ▶ +1 full throttle forward,
 - ▶ -1 full throttle backwards,
 - ▶ 0 zero throttle.
- ▶ The system dynamics are as below

$$\begin{aligned}x_{t+1} &\leftarrow x + \dot{x}_{t+1} \\ \dot{x}_{t+1} &\leftarrow \dot{x}_t + 0.001A_t - 0.025\cos(3x_t),\end{aligned}$$

where x_t denotes the position and \dot{x}_t the velocity of the car.

Example: Mountain Car

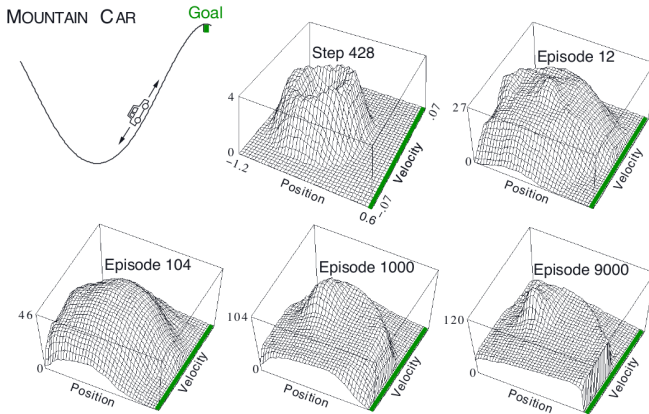


Figure: R. Sutton, A. Barto, MIT Press, 2017

Example: Mountain Car

- ▶ Mountain car is a standard application for delayed reward:
Driving towards the exit point is not the right way.
- ▶ Step 428 has a symmetric shape, because all initially visited states are valued worse than the default value unexplored states.
- ▶ Consequently, the agent decides to explore for long episodes even though $\epsilon = 0$.

Average reward

- ▶ Applicable to continuing problems.
- ▶ Discounting causes problems with function approximation.
- ▶ The goodness measure for π is average reward / time step

$$\begin{aligned}\eta(\pi) &\triangleq \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \mathbb{E}[R_t | A_{0:t-1} \sim \pi] \\ &= \lim_{t \rightarrow \infty} \mathbb{E}[R_t | A_{0:t-1} \sim \pi] \\ &= \sum_s d_\pi(s) \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) r,\end{aligned}$$

where

$$d_\pi(s) \triangleq \lim_{t \rightarrow \infty} P[S_t = s | A_{0:t-1} \sim \pi]$$

is the steady-state distribution.

Ergodicity

- ▶ The steady-state distribution is independent of the starting state S_0 .
- ▶ Markov chains with a steady-state distribution are called *ergodic*.
- ▶ Taking an action wrt π keeps $d_\pi(s)$ unchanged:

$$\sum_s d_\pi(s) \sum_a \pi(a|s, \mathbf{w}) p(s'|s, a) = d_\pi(s')$$

Differential return

$$G_t \triangleq R_{t+1} - \eta(\pi) + R_{t+2} - \eta(\pi) + \dots$$

Mind that there is *no discount*! The related value function is also named as the **differential value function**

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{r,s'} p(s', r|s, a) \left[r - \eta(\pi) + v_{\pi}(s') \right]$$

$$q_{\pi}(s, a) = \sum_{r,s'} p(s', r|s, a) \left[r - \eta(\pi) + \sum_{a'} \pi(a'|s') q_{\pi}(s', a') \right]$$

$$v_*(s) = \max_a \sum_{r,s'} p(s', r|s, a) \left[r - \eta(\pi) + q_*(s', a) \right]$$

$$q_*(s, a) = \sum_{r,s'} p(s', r|s, a) \left[r - \eta(\pi) + \max_{a'} q_*(s', a') \right]$$

Differential TD error

For state-value function

$$\delta_t \triangleq R_{t+1} - \bar{R}_t + \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})$$

For action-value function

$$\delta_t \triangleq R_{t+1} - \bar{R}_t + \hat{q}(S_{t+1}, A_t, \mathbf{w}) - \hat{q}(S_t, A_t, \mathbf{w})$$

Differential semi-gradient Sarsa for control

Input: a differentiable $\hat{q} : \mathcal{S} \times \mathcal{A} \times \mathbb{R}^d \rightarrow \mathbb{R}$, step sizes $\alpha, \beta > 0$

Initialize $\mathbf{w} \in \mathbb{R}^d$, \bar{R} , S , A

repeat (for each step)

Take action A , observe R, S'

Choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g. ϵ -greedy)

$\delta \leftarrow R - \bar{R} + \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w})$

$\bar{R} \leftarrow \bar{R} + \beta \delta$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha \delta \nabla \hat{q}(S, A, \mathbf{w})$

$S \leftarrow S'; A \leftarrow A'$

Futility of the discount factor

$$\begin{aligned} J(\pi) &= \sum_s d_\pi(s) v_\pi^\gamma(s) \quad (v_\pi^\gamma(s) \text{ is the discounted value function}) \\ &= \sum_s d_\pi(s) \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) [r + \gamma v_\pi^\gamma(s')] \\ &= \eta(\pi) + \sum_s d_\pi(s) \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) \gamma v_\pi^\gamma(s') \\ &= \eta(\pi) + \gamma \sum_{s'} v_\pi^\gamma(s') \sum_s d_\pi(s) \sum_a \pi(a|s) p(s'|s,a) \\ &= \eta(\pi) + \gamma \sum_{s'} v_\pi^\gamma(s') d_\pi(s') \\ &= \eta(\pi) + \gamma J(\pi) = \eta(\pi) + \gamma \eta(\pi) + \gamma^2 J(\pi) \\ &= \eta(\pi) + \gamma \eta(\pi) + \gamma^2 \eta(\pi) + \gamma^3 J(\pi) + \dots = \frac{1}{1-\gamma} \eta(\pi) \end{aligned}$$

Futility of the discount factor

- ▶ As $J(\pi) = \frac{1}{1-\gamma} \eta(\pi)$, the discount factor scales all policies with the same factor $1/(1-\gamma)$.
- ▶ Hence, does not affect their *ordering*!
- ▶ Discounting still matters in the episodic case!